

Building Reproducible and Replicable Projects

R Workflows

Daniel Chen @chendaniely

NYR Conference 2019


https://github.com/chendaniely/rstatsnyc_2019-workflow

hi!

I'm Daniel



- PhD Student: Virginia Tech
- Instructor
- Data Scientist
- Community Member
- RStudio Intern!
- Author:





D@niel Chen @pycon

@chendaniely




Are there any [#rstats](#) users at [#PyCon19](#)??
My next stop is [#rstatsnyc](#)!

[#pycon](#) [#nycdatamafia](#)



10:18 AM - 5 May 2019 from [Pycon 2019](#)

1 Retweet 9 Likes



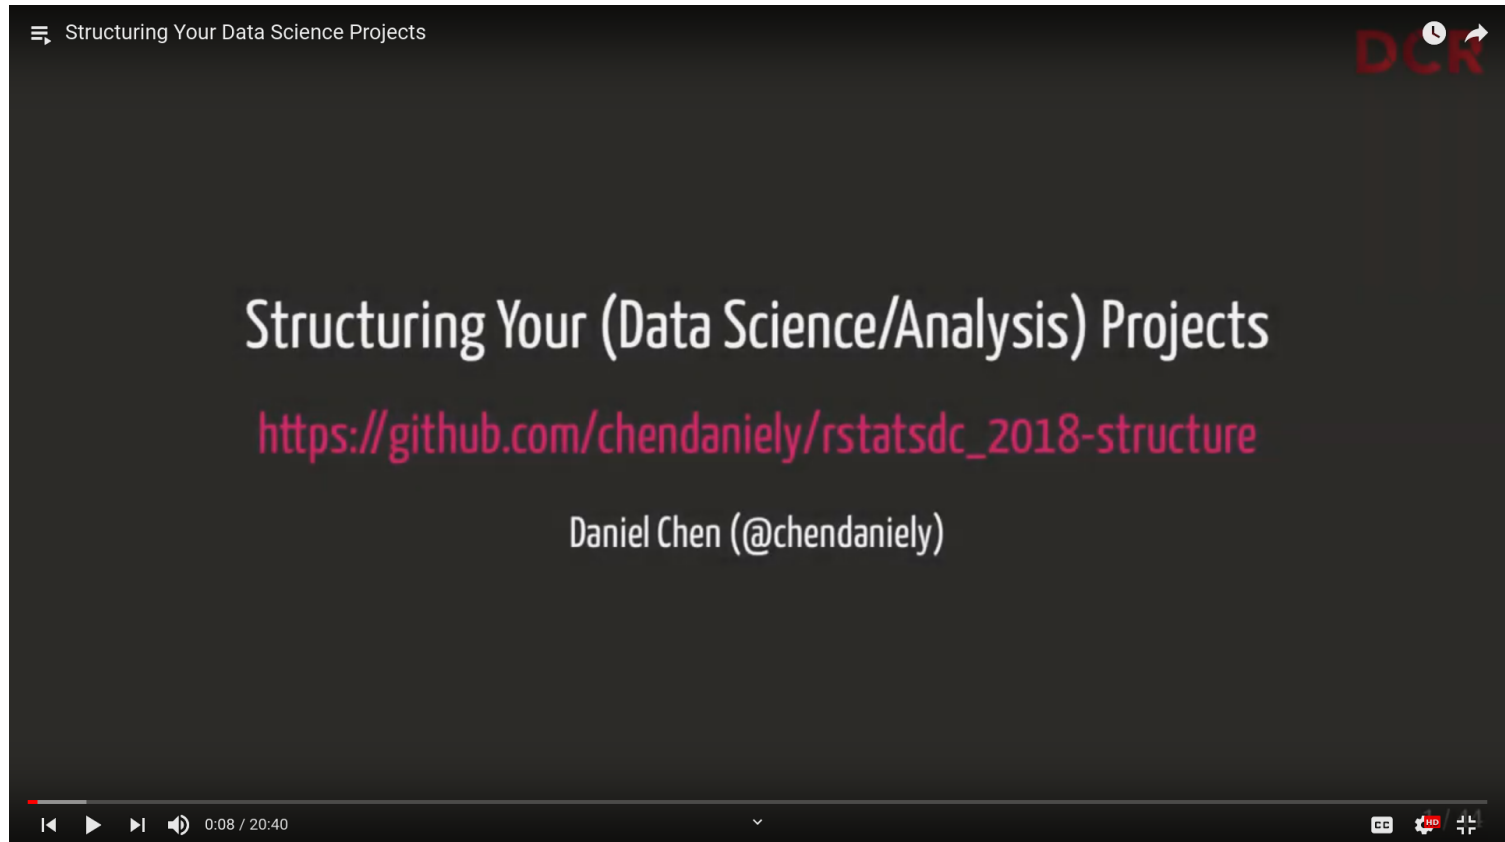
Caves!

Ginnie Springs, FL



- <https://www.submergedscuba.com/>
- <http://www.scuba.org.vt.edu/>

Continuing the DCR 2018 Story



<https://youtu.be/UQH38s3DyA>

Previously...

Structuring Your Data Science Projects

We are happy when our code just runs

R has given us the tools to make your projects more structured and organized

Many people converge on very similar project templates

It doesn't matter where you are in your learning path

tl;dr

- I just want stuff to run the first time around

In sum...

1. Use R
2. Make a project
3. Organize the project into folders and use `here::here()` to get project relative paths
4. Break up scripts into smaller pieces
5. RMarkdown for things you want to show
6. Put functions in R so your analysis is package ready and write `Makefiles`, shell scripts, or other build scripts and link your projects to scholarship so your figures and tables are always up to date

Use R and put all your scripts in one place


Create a project so you don't have to manually set working directories

Organize your R project into folders so you don't end up with a huge list of files to hunt down

Take the scripts in your project and break them up into smaller scripts that do smaller tasks. No more scroll through 100s of line of code to re-render 1 plot

Create R scripts to do the data processing and use Rmd files to coherently show your results. No more fumbling through stuff nobody cares about at meetings and you'll look like a badass

Create a Makefile to re-run your entire analysis pipeline, get updated figures. Put your functions in a "R" folder so it is package ready. Create a separate project for your full scholarly report, link the analysis code into the report, and have a report/presentation that always has the latest figures/tables/data



Where we originally started

1. A (billboard) dataset
2. A script that had everything in it

https://github.com/chendaniely/rstatsdc_2018-structure/blob/master/01-just_starting_out/analysis.R

Reproducibility vs Replicability

- **reproducibility** - the extent to which consistent results are obtained when an experiment is repeated
- **replicability** - the ability of a scientific experiment or trial to be repeated to obtain a consistent result

```
☺ find . -type f -not -path './.Rproj*'
./analysis/billboard_eda/03-eda.Rmd
./analysis/billboard_eda/02-01-clean.R
./analysis/billboard_eda/02-02-tidy.R
./analysis/billboard_eda/02-03-normalize.R
./analysis/billboard_eda/01-load.R
./analysis/billboard_eda/04-model.Rmd
./Makefile
./output/billboard_reports/03-eda.html
./output/billboard_reports/04-model.html
./output/billboard_model_plots/coefs_predict_rank_week.png
./output/billboard_model_plots/coefs_predict_rank_week_no_intercept.png
./output/billboard_model_plots/coefs_predict_rank_week_artist_sorted.png
./output/billboard_model_plots/coefs_predict_rank_week_artist.png
./output/billboard_rank_plots/avg_rnk_by_week.png
./output/billboard_rank_plots/avg_rnk_by_month.png
./output/billboard_rank_plots/avg_rank_by_week_across_months.png
./data/original/billboard.csv
./data/processed/billboard/billboard_clean.csv
./data/processed/billboard/billboard.csv
./data/processed/billboard/.gitkeep
./data/processed/billboard/songs.csv
./data/processed/billboard/billboard_tidy.csv
./data/processed/billboard/rank.csv
./ex-01-make.Rproj
```

Running scripts from command line

- Running from the root project directory

```
$ Rscript ./analysis/billboard_eda/01-load.R
```

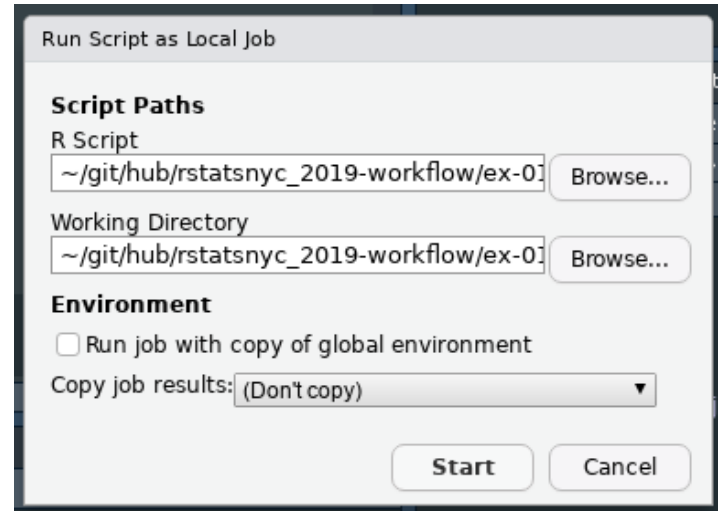

- Running the script with here and rprojroot libraries

```
$ cd ./analysis/billboard_eda  
$ Rscript 01-load.R
```

- Background the script

```
$ nohup Rscript 01-load.R &
```

Jobs tab in RStudio!



Build scripts (Make)

A task consists of:

```
target: dependency1 dependency2  
    action1  
    action2
```

Build scripts (Make)

```
BILLBOARD=./analysis/billboard_eda/

all : commands

## commands          : show all commands.
commands :
    @grep -E '^##' Makefile | sed -e 's/## //g'

## billboard_eda    : re-generate billboard eda analysis
billboard_eda :
    Rscript ${BILLBOARD}/01*
    Rscript ${BILLBOARD}/02-01*
    Rscript ${BILLBOARD}/02-02*
    Rscript ${BILLBOARD}/02-03*
    Rscript -e "rmarkdown::render(here::here('./analysis/billboard_eda/02-01.Rmd'), output_file='./analysis/billboard_eda/02-01.html', quiet=TRUE)"
    Rscript -e "rmarkdown::render(here::here('./analysis/billboard_eda/02-02.Rmd'), output_file='./analysis/billboard_eda/02-02.html', quiet=TRUE)"
    Rscript -e "rmarkdown::render(here::here('./analysis/billboard_eda/02-03.Rmd'), output_file='./analysis/billboard_eda/02-03.html', quiet=TRUE)"

## clean            : clean up junk files.
clean :
    find data/processed/ -type f -name '*.csv' | xargs rm
    find analysis/ type f -name '*.html' | xargs rm
```


File inputs

file	input
01-load.R	'./data/original/billboard.csv'
02-01-clean.R	'./data/processed/billboard/billboard.csv'
02-02-tidy.R	'./data/processed/billboard/billboard_clean.csv'
02-03-normalize.R	'./data/processed/billboard/billboard_tidy.csv'
03-eda.Rmd	'./data/processed/billboard/rank.csv'
04-model.Rmd	'./data/processed/billboard/billboard_tidy.csv'

File outputs

file	output
01-load.R	'./data/processed/billboard/billboard.csv'
02-01-clean.R	'./data/processed/billboard/billboard_clean.csv'
02-02-tidy.R	'./data/processed/billboard/billboard_tidy.csv'
02-03-normalize.R	'./data/processed/billboard//songs.csv'; ./data/processed/billboard/rank.cs
03-eda.Rmd	'./output/billboard_rank_plots/avg_rnk_by_week.png'; ./output/billboard_rank_plots/avg_rnk_by_month.png; ./output/billboard_rank_plots/avg_rank_by_week_across_months.png
04-model.Rmd	'./output/billboard_model_plots/coefs_predict_rank_week.png'; ./output/billboard_model_plots/coefs_predict_rank_week_no_intercept.pr ./output/billboard_model_plots/coefs_predict_rank_week_artist.png

Draw your build scripts (Make)

Build scripts (Make) -- "better" Variables

```
BILLBOARD=./analysis/billboard_eda  
DO=./data/original  
DP_BILLBOARD=./data/processed/billboard  
O_BILLBOARD_PLOTS=./output/billboard_rank_plots  
O_BILLBOARD_REPORTS=./output/billboard_reports
```

Build scripts (Make) -- "better" tidy

```
## tidy          : make tidy dataset for analysis
.PHONY: tidy
tidy: $(DP_BILLBOARD)/billboard_tidy.csv \
      $(DP_BILLBOARD)/songs.csv \
      $(DP_BILLBOARD)/rank.csv

$(DP_BILLBOARD)/billboard.csv: $(DO)/billboard.csv $(BILLBOARD)/01-load.R
Rscript $(BILLBOARD)/01-load.R
$(DP_BILLBOARD)/billboard_clean.csv: $(DP_BILLBOARD)/billboard.csv \
                                      $(BILLBOARD)/02-01*.R
Rscript $(BILLBOARD)/02-01*.R
$(DP_BILLBOARD)/billboard_tidy.csv: $(DP_BILLBOARD)/billboard_clean.csv \
                                    $(BILLBOARD)/02-02*.R
Rscript $(BILLBOARD)/02-02*.R
$(DP_BILLBOARD)/songs.csv: $(DP_BILLBOARD)/billboard_tidy.csv \
                           $(BILLBOARD)/02-03*.R
Rscript $(BILLBOARD)/02-03*.R
$(DP_BILLBOARD)/rank.csv: $(DP_BILLBOARD)/billboard_tidy.csv \
                          $(BILLBOARD)/02-03*.R
Rscript $(BILLBOARD)/02-03*.R
```

Build scripts (Make) -- "better" eda and model

```
## eda                                : create the eda report
.PHONY: eda
eda: $(DP_BILLBOARD)/rank.csv \
     $(O_BILLBOARD_REPORTS)/03-eda.html

$(O_BILLBOARD_REPORTS)/03-eda.html: $(BILLBOARD)/03-eda.Rmd
    Rscript -e "rmarkdown::render(here::here('./analysis/billboard_ec

## model                              : create the model report
.PHONY: model
model: $(DP_BILLBOARD)/billboard_tidy.csv \
      $(O_BILLBOARD_REPORTS)/04-model.html
```

Build scripts (Make) -- "better" eda and model

```
## eda                                : create the eda report
.PHONY: eda
eda: $(DP_BILLBOARD)/rank.csv \
     $(O_BILLBOARD_REPORTS)/03-eda.html

$(O_BILLBOARD_REPORTS)/03-eda.html: $(BILLBOARD)/03-eda.Rmd
    Rscript -e "rmarkdown::render(here::here('./analysis/billboard_ec

## model                               : create the model report
.PHONY: model
model: $(DP_BILLBOARD)/billboard_tidy.csv \
      $(O_BILLBOARD_REPORTS)/04-model.html
```


Build scripts (Make) -- "better" reports

```
## reports          : create the eda and model reports
.PHONY: reports
reports: $(O_BILLBOARD_REPORTS)/03-eda.html \
        $(O_BILLBOARD_REPORTS)/04-model.html

## tidy_and_reports: run the entire tidy and report pipeline
.PHONY: tidy_and_reports
tidy_and_reports: tidy reports
```

```
make -f Makefile2 clean
make -f Makefile2 tidy_and_reports
# make -f Makefile2 tidy
# make -f Makefile2 reports
```

Introducing ... drake!

```
# Install the latest stable release from CRAN.  
install.packages("drake")
```

- Just like make it will only run scripts for things that are out-of-date.
- If something does not need to be re-run, it won't re-run it.
- The manual: <https://ropenscilabs.github.io/drake-manual>
- The docs: <https://ropensci.github.io/drake/>

But... It's not exactly like make

- drake plan: data frame with columns named target and command
- target: R object
- command: expression to produce it

A drake plan

```
plan <- drake_plan(  
  raw_data = readxl::read_excel(file_in("raw_data.xlsx")),  
  data = raw_data %>%  
    mutate(Species = forcats::fct_inorder(Species)),  
  hist = create_plot(data),  
  fit = lm(Sepal.Width ~ Petal.Width + Species, data),  
  report = rmarkdown::render(  
    knitr_in("report.Rmd"),  
    output_file = file_out("report.html"),  
    quiet = TRUE  
  )  
)  
plan  
#> # A tibble: 5 x 2  
#>   target      command  
#>   <chr>      <expr>  
#> 1 raw_data  readxl::read_excel(file_in("raw_data.xlsx"))  
#> 2 data      raw_data %>% mutate(Species = forcats::fct_inorder(Spec  
#> 3 hist      create_plot(data)  
#> 4 fit       lm(Sepal.Width ~ Petal.Width + Species, data)  
#> 5 report    rmarkdown::render(knitr_in("report.Rmd"), output_file :
```

What about all my scripts?

drake plans rely more on in-memory functions

```
bad_plan <- drake_plan(  
  my_data = source(file_in("get_data.R")),  
  my_analysis = source(file_in("analyze_data.R")),  
  my_summaries = source(file_in("summarize_data.R"))  
)
```

I personally can't use this

- I'd have to re-write all my scripts
- Patch really bad code
 - wrapping each of my scripts around a function call
 - still need to point out my inputs and outputs

Instead of just complaining...

- Project I'm working on this summer (RStudio) -- grader
- Automatically grade student's code given solution code and report why/where the code is wrong
- Given a script, what are the file inputs and outputs
 - More than just string extration

The plan

- Use Abstract Syntax Trees (ASTs)

```
library(lobstr)
```

```
ast(a + b + c)
```

```
#>   
#>   
#>   
#>   
#> 
```

- Use implement my find input/output script dectector
- Make my workflow work in drake?

Thanks!

@chendaniely

Slides: https://github.com/chendaniely/rstatsnyc_2019-workflow